

VICTORIA UNIVERSITY OF WELLINGTON
Te Whare Wānanga o te Ūpoko o te Ika a Māui



School of Engineering and Computer Science
Te Kura Mātai Pūkaha, Pūrorohiko

PO Box 600
Wellington
New Zealand

Tel: +64 4 463 5341
Fax: +64 4 463 5045
Internet: office@ecs.vuw.ac.nz

QuickCheck for Whiley
Preliminary Report

Janice Chin

Supervisors: David Pearce and Lindsay Groves

Submitted in partial fulfilment of the requirements for
Bachelor of Engineering with Honours.

Abstract

Contents

1	Introduction	1
2	Background Survey	2
2.1	Whiley	2
2.2	QuickCheck	2
2.3	Concolic Unit Testing Engine (CUTE)	3
2.4	Randoop	3
3	Work Completed	4
3.1	Test Case Generation Techniques	4
3.1.1	Random Test Case Generation	4
3.1.2	Exhaustive Test Case Generation	4
3.2	Data types generated	4
3.3	Integer Range Analysis	5
3.4	Evaluating the tool	6
4	Future Plan	7
5	Request for Feedback	8
A	Project Proposal	10

Chapter 1

Introduction

The detection and elimination of bugs is important in software development as it improves the quality and reliability of software.

To help eliminate errors in code, the Whiley programming language was developed to verify code using formal specifications. Ideally, programs written in Whiley should be error free. To achieve this goal, Whiley contains a verifying compiler which employs the use of specifications written by a developer to check for common errors such as accessing an index of an array which is outside its boundaries [1].

Listing 1.1 illustrates an example of a Whiley program with a function, `min()` which finds the minimum value of two integers. Executing `min(2, 3)` will return the integer, 2.

Listing 1.1: Whiley program for the min function

```
1 function min(int x, int y) -> (int r)
2 ensures r == x ==> x <= y
3 ensures r == y ==> y <= x:
4 if x <= y:
5     return x
6 else:
7     return y
```

Currently, the verifying compiler has limitations when evaluating complex pre- and post-conditions. For example in Listing 1.2, the post-condition is falsely identified as not holding by the verifying compiler even though the program does meet the post-condition.

Listing 1.2: Whiley program for the square function

```
1 function square(int x) -> (int r)
2     ensures r >= 0:
3     return x*x
```

Consequently, tests would need to be created to detect if there are further problems within the program. However, writing and running tests can be tedious and costly. Furthermore, it is difficult to write tests for all possible cases and be able to detect all bugs. Instead of manually creating tests, a tool could be created that automatically generates and executes tests on a program such as QuickCheck which was implemented in Haskell by Koen Claessen and John Hughes [2].

This project aims to implement an automated test-case generator for the Whiley programming language, based on the QuickCheck tool. As a result, the automatic test-case generator will improve software quality and increase confidence in unverifiable code.

Chapter 2

Background Survey

There exists a number of different test-case generation tools for different languages such as QuickCheck for Haskell, Quviq QuickCheck for Erlang, Concolic Unit Testing Engine (CUTE) for C and Randoop for Java and .NET.

QuickCheck for Whiley is specifically created for the Whiley programming language. Another key difference while making QuickCheck for Whiley is that different techniques are used to generate test data instead of just one technique.

2.1 Whiley

Whiley is a hybrid imperative and functional programming language developed by David Pearce.

A notable aspect about Whiley is the verifying compiler used in conjunction with explicit specifications to verify programs are correct.

Whiley files are compiled with or without static verification into a WyIL file which is a binary form of the Whiley program in an intermediate language.

2.2 QuickCheck

QuickCheck is a tool implemented by Koen Claessen and John Hughes [2] to test Haskell programs.

To check if a program is correct, QuickCheck uses property-based testing which uses conditions that will always hold true. Users define the conditions as Haskell functions in terms of formal specifications to validate functions under test.

QuickCheck then generates a large number of test cases automatically using random testing. Input values are randomly generated for each test (within a range) using generators that correspond to the input value's type. QuickCheck contains generators for most of Haskell's predefined types and for functions. User-defined types require the tester to specify their own custom generators with a bounded size if the type is recursive such as a binary tree. [2].

A problem with random testing is the distribution of test data. Ideally, the distribution should adhere to the distribution of actual data which is often uniformly distributed. For example, executing tests for Listing 2.2 could be skewed towards test data with short lists. The distribution of data is controlled by the tester by creating a custom generator for the data type with weighted frequencies on the methods used to generate the data.

An example of using QuickCheck would be reversing a list.

Firstly, a property is defined in Listing 2.1 where the list `xs`, should be the same as reversing `xs` twice.

Listing 2.1: Property for reversing a list in QuickCheck

```
propReverseTwice xs = reverse (reverse xs) == xs
```

QuickCheck is then executed by importing the property and passing it into the interpreter as shown in Listing 2.2

Listing 2.2: Executing tests to check a list is reversed

```
Main:> quickCheck propReverseTwice
Ok: passed 100 tests .
```

This will create a large number of test cases by using the user-defined property and random input values to check the property holds. QuickCheck reports various test statistics including the number of tests that have passed or failed [2].

2.3 Concolic Unit Testing Engine (CUTE)

Concolic Unit Testing Engine (CUTE) generates test inputs using a combination of symbolic and concrete execution [3]. The aim of CUTE is to provide input values to explore all feasible execution paths of a program and thus achieve high path coverage [3]. CUTE has been applied to C programs [3].

Random testing may generate inputs with the same behaviour leading to its redundancy [3]. Furthermore, the probability of selecting inputs that will detect errors using random testing is small [3]. Therefore, CUTE incrementally generates concrete inputs using symbolic execution to discover feasible paths [3].

Firstly, CUTE executes the program with arbitrary inputs using concrete and symbolic execution concurrently [3]. In concrete execution, the program is executed normally with the input values [3]. At the same time, symbolic execution is run through the same path. Using the symbolic variables, constraints from branching expressions are discovered and stored [3]. The last constraint applied is negated so that the next test run will explore a different feasible path [3]. The constraint is solved to limit and determine the possible inputs to execute in the next test [3].

2.4 Randoop

Randoop generates unit tests using feedback-directed random test generation [4], [5]. Method sequences from previous tests are used to help generate subsequent tests [4], [5]. As a result, a test suite is outputted with successful and unsuccessful tests [4], [5].

Randoop tests classes by executing a sequence of methods as a test [5]. Method sequences are created incrementally by randomly selecting method calls and using arguments from previous sequences [5]. The sequence is then executed and checked against contracts [4]. Contracts are built into Randoop or optionally defined by the user [4]. If a sequence breaks a contract then the test is outputted as a contract-violating test [4]. If sequence was successful, then the test is outputted as a regression test [4]. Successful sequences that can be extended or are not redundant are then used in subsequent tests [4].

Chapter 3

Work Completed

An implementation of the tool has been created with generation of core and more complex data types. The tool has also been extended to use integer range analysis to constrain input values generated.

For a user to use the tool, a user must pass in a WyIL file and specify the test case generation technique to use, number of tests they wish to generate and ranges for generating integer values.

The Whiley Theorem Prover is used to validate inputs and outputs from tests. Valid inputs based on verifying pre-conditions and constraints on the input type are used. A successful test is checked by using the function's post-condition and constraints on the output type.

Currently, only functions within the program with a post-condition are able to be tested as they can be verified.

3.1 Test Case Generation Techniques

QuickCheck for Whiley currently employs random test-case and exhaustive test case generation to generate test cases:

3.1.1 Random Test Case Generation

A fixed number of tests determined by the user is randomly generated. Previous inputs are not considered which may lead to the same inputs being generated. Algorithm ... was considered to generate unique inputs but has not been implemented yet.

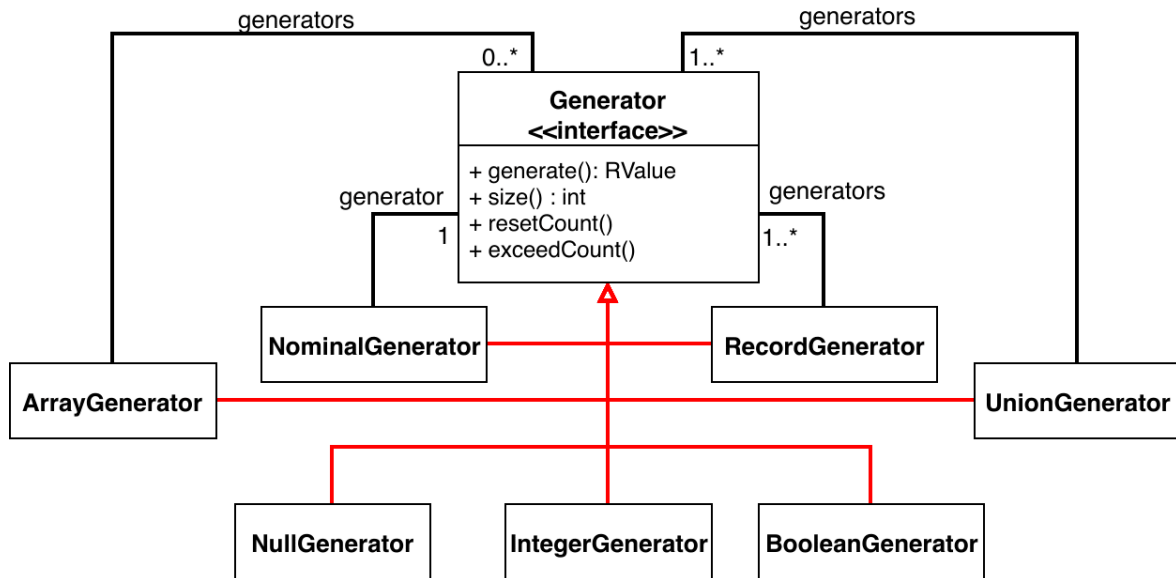
3.1.2 Exhaustive Test Case Generation

A fixed number of tests determined by the user are exhaustively generated within a bounded range from the minimum bound to the maximum bound. All possible input combinations for a function are tested starting from the combination generated from the minimum bound.

3.2 Data types generated

Different generators are created for the different types defined in Whiley. Figure 3.1 illustrates the structure of the generators used. Each generator generates a value. It also has a size for the number of input combinations and `resetCount()` and `exceedCount()` methods for monitoring input combinations generated for test-case generation.

Figure 3.1: Class diagram of Generators in QuickCheck for Whiley



Null Generate a null value.

Boolean Generates a boolean which is either true or false.

Integer Generates an integer between a bounded range. This range may be modified based on constraints applied to the integer.

Array Generates an array with a specific element type between a bounded size range. The size of arrays are currently limited to a maximum size of three elements due to the performance cost of generating larger arrays.

Nominal A nominal represents a user-defined type by redefining a type with a different name [1]. Therefore, a generator for a nominal wraps a generator for a different type to generate a value. A nominal can also have a type constraint/invariant applied to it [1].

Record "A record type describes the set of all compound values made from one or more fields, each of which has a unique name and a corresponding type [1]." A closed record is generated by generating values that correspond to each field. Generators corresponding to each field is required as each field may have different types.

Union A union is made up of multiple types where the value can correspond to any one of the types declared. For example, the type `bool | int` can hold either a boolean or an integer value. A union is generated by generating a value from one of the declared types. A generator corresponding to each type is required.

Values for the union type could be skewed towards generating values for only type. It was important that the different types for each union was fairly generated. Therefore, the union generator was implemented so that a value from a different type would be generated each time by iterating through the generators within the union generator.

3.3 Integer Range Analysis

Integer range analysis is ...

3.4 Evaluating the tool

The tool has been evaluated by executing a variety of tests notably, the test suite for the Whiley Compiler.

Chapter 4

Future Plan

Chapter 5

Request for Feedback

Bibliography

- [1] D. Pearce, “The Whyley Language Specification,” 2014.
- [2] K. Claessen and J. Hughes, “QuickCheck: A Lightweight Tool for Random Testing of Haskell Programs,” *SIGPLAN Not.*, vol. 46, pp. 53–64, May 2011.
- [3] K. Sen, D. Marinov, and G. Agha, “CUTE: A Concolic Unit Testing Engine for C,” *SIGSOFT Softw. Eng. Notes*, vol. 30, pp. 263–272, Sept. 2005.
- [4] C. Pacheco, S. K. Lahiri, M. D. Ernst, and T. Ball, “Feedback-directed random test generation,” in *Proceedings of the 29th International Conference on Software Engineering, ICSE '07*, (Washington, DC, USA), pp. 75–84, IEEE Computer Society, 2007.
- [5] C. Pacheco and M. D. Ernst, “Randoop: Feedback-directed random testing for java,” in *Companion to the 22Nd ACM SIGPLAN Conference on Object-oriented Programming Systems and Applications Companion, OOPSLA '07*, (New York, NY, USA), pp. 815–816, ACM, 2007.

Appendix A

Project Proposal

VICTORIA UNIVERSITY OF WELLINGTON
Te Whare Wānanga o te Ūpoko o te Ika a Māui



School of Engineering and Computer Science
Te Kura Mātai Pūkaha, Pūrorohiko

PO Box 600
Wellington
New Zealand

Tel: +64 4 463 5341
Fax: +64 4 463 5045
Internet: office@ecs.vuw.ac.nz

QuickCheck for Whiley

Janice Chin

Supervisors: David Pearce and Lindsay Groves

Submitted in partial fulfilment of the requirements for
Bachelor of Engineering with Honours.

Abstract

This document is a project proposal for the project, QuickCheck for Whiley. The project aims to improve software quality of Whiley programs by implementing an automated test-case generator for Whiley. The solution will be evaluated by checking if it can detect bugs in a small benchmark set. No special resources are required for this project.

1. Introduction

Testing is an important process in software development as it helps detect the presence of bugs. However, writing and running tests can be tedious and costly. Furthermore, it is difficult to write tests for all possible cases and be able to detect all bugs.

An automated test-case generator called QuickCheck was implemented in Haskell by Koen Claessen and John Hughes [1] which alleviates these issues.

This lightweight tool employs two testing ideas. Firstly, QuickCheck uses property-based testing which uses conditions that will always hold true. Users define the conditions in terms of formal specifications to validate functions under test.

Secondly, QuickCheck generates test cases automatically using random testing. Input values are randomly generated for each test (within a range) using generators that correspond to the input value's type. QuickCheck contains generators for most of Haskell's pre-defined types and for functions. It requires the developer to specify their own generators for user-defined types [1].

An example of using QuickCheck would be reversing a list.

Firstly, a property is defined where the list `xs`, should be the same as reversing `xs` twice.

```
propReverseTwice xs = reverse (reverse xs) == xs
```

QuickCheck is then executed by importing the property and passing into the interpreter:

```
Main:> quickCheck propReverseTwice
Ok: passed 100 tests.
```

This will create a large number of test cases by using the user-defined property and random input values to check the property holds. QuickCheck reports various test statistics including the number of tests that have passed or failed [1].

QuickCheck has been re-implemented in other languages including Scala, Java and C++. A commercial version of QuickCheck in Erlang called Quviq QuickCheck is co-founded by one of the original developers, John Hughes and is an extension of the original QuickCheck for Haskell [2].

This project is about implementing an automated test-case generator for the programming language, Whiley based on the QuickCheck tool.

2. The Problem

Whiley is a programming language, developed to verify code and eliminate errors using formal specifications. Ideally, programs written in Whiley should be error free. To achieve this goal, Whiley contains a verifying compiler which employs the use of specifications written by a developer to check for common errors such as accessing an index of an array which is outside its boundaries [3].

Below is an example of a Whiley program with a function, `min()` which finds the minimum value of two integers. This program prints out "The smallest integer is 3" when executed.

Listing 1: Whiley program for the min function

```
1 import std::io
2 import std::ascii
3
4 method main():
5     int x = 5
6     int y = 3
7     int smallest = min(x, y)
8     io::print("The smallest integer is ")
9     io::println(smallest)
10
11 function min(int x, int y) -> (int r)
12 ensures r == x ==> x <= y
13 ensures r == y ==> y <= x:
14     if x <= y:
15         return x
16     else:
17         return y
```

Currently, the verifying compiler has limitations when evaluating complex pre- and post-conditions. For example, a post-condition could be falsely identified as not holding by the verifying compiler even though the program does meet the post-condition. Therefore, this project aims to implement an automated test-case generator in Whiley to improve software quality and increase confidence in unverifiable code.

3. Proposed Solution

The automated test-case generator needs to read a Whiley program and then generate tests for functions in the program. Input values that adhere to the function's preconditions will be used in testing. Strategies for generating input values include random generation and dynamic generation. The success of each test is determined by using the function's postconditions. The test results should show the percentage of tests that passed and the inputs used in the failed tests.

An example of a function for test-case generation would be for the min function introduced in Listing 1.

Possible input values automatically generated include:

- min(4, 4) would result in r = 4
- min(6, 7) would result in r = 7
- min(10, 5) would result in r = 10
- min(-2, 8) would result in r = 8

Several components developed for the Whiley language will need to be used for the automated test-case generator. The Whiley interpreter will be used for checking candidate values meet precondition and validating test outputs meet the postconditions. It will also be used for analysing the program structure such as the names of functions declared and their parameters. The Whiley Compiler will be used for generating input data for different types.

After basic test-case generation is completed, more complex types can be generated for testing and/or further enhancements to the tool can be made.

Possible enhancements include:

- Weight on the frequency of test values
- Shrinking failed test cases
- Classifying tests
- Using dynamic symbolic execution

Regular meetings (weekly or fortnightly) with the primary supervisor, David Pearce will be held for assistance and guidance in the project.

Timeline

A gantt chart of the project can be found on the project's repository.

Trimester 1

Output	Estimated Time	Start Date	Complete by
Produce bibliography	2 weeks	19/3/18	9/4/18 (Week 5)
Produce project proposal	2 weeks	19/3/18	9/4/18 (Week 5)
Implement an automated test-case generator for Whiley programs for the primitive types: bool, byte, int, null	3 weeks	19/3/18	16/4/18 (Week 6)
Implement automated test-case generation for more complex types: void, array, union and records	3 weeks	16/4/18	14/5/18 (Week 9)
Produce preliminary report	3 weeks	14/5/18	10/6/18 (Week 12)
Extend the automatic test generator to be able to generate other types or use methods for better test case distribution (weighting, classification)	3 weeks	11/6/18	2/7/18 (Exam period)

Trimester 2

Output	Estimated Time	Start Date	Complete by
Produce slides for presentation of preliminary report	1 week	2/7/18	9/7/18 (Break)
Extend the automatic test generator to include other methods of testing (symbolic) or be able to generate other types	3 weeks	9/7/18	30/7/18 (Week 3)
Produce draft of final report	6 weeks	30/7/18	15/9/18 (Week 7)
Complete implementation of automated test generator	3 weeks	15/9/18	5/10/18 (Week 10)
Finalise final report	2 weeks	6/10/18	21/10/18 (Week 12)
Prepare slides for final presentation	3 weeks	22/10/18	16/11/18 (Exam period)

4. Evaluating your Solution

This project will be evaluated by checking if the developed tool can detect bugs that have been inserted into a small benchmark set.

5. Resource Requirements

No special resources are required. Only the use of the ECS laboratories is required.

Bibliography

- [1] K. Claessen and J. Hughes, “Quickcheck: A lightweight tool for random testing of haskell programs,” *SIGPLAN Not.*, vol. 46, pp. 53–64, May 2011.
- [2] J. Hughes, “Quickcheck testing for fun and profit,” in *Proceedings of the 9th International Conference on Practical Aspects of Declarative Languages*, PADL’07, (Berlin, Heidelberg), pp. 1–32, Springer-Verlag, 2007.
- [3] D. Pearce, “The Whyley Language Specification,” 2014.